

# Text Analytics - Amazon Fine Food Reviews

AMRITA SHARMA &  
PREETHI RANGANATHAN

# Table of Contents

- Data Overview
- Data Cleanup
- Business Questions
- Dataset Preparation
- Features
- Model Analysis
- Insights

# Dataset Overview

Dataset consists of 568,452 Amazon Fine Food Reviews.

Attributes:

- Review id
- Product Category Id
- User Id
- User profile name
- Number of users who voted for the helpfulness survey
- Number of users who found the review helpful
- Rating (1-5)
- Review Date & Time
- Review
- Review Title

# Some sample reviews

## Positive review

Text

1:

I have bought several of the Vitality canned dog food products and have found them all to be of good quality. The product looks more like a stew than a processed meat and it smells better. My Labrador is finicky and she appreciates this product better than most.

## Negative Review

Definitely not worth buying flavored water with a few teaspoons of beans and rice that doesn't taste like normal beans and rice. I won't ever buy this again!

## Neutral Review

It's great to have agave in a portable format. But it is difficult to open. The directions say to pinch to open. They are not so easy to pinch. I have found if you bend the tube about 1 inch above the end then pinch the end it helps. But the agave often gets on your fingers.

# Data Cleanup - Spam Review Identification

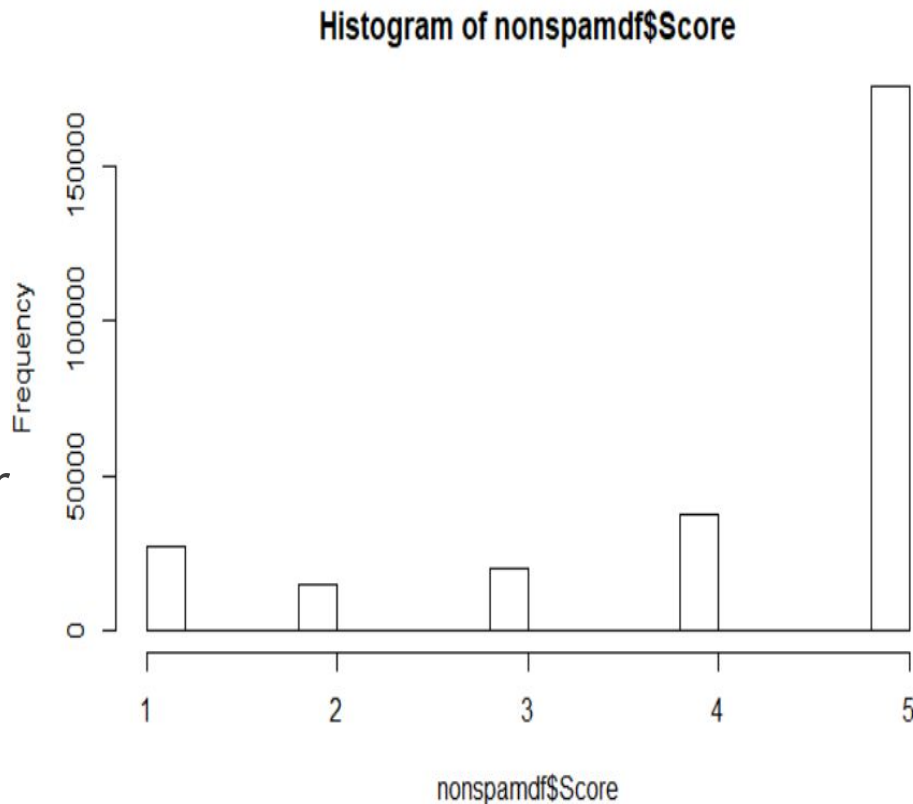
- Fake reviews
- To manipulate online customers' opinions on products being sold
- Usually provided by bots
- To make their own business look good or damage someone else's business
- Multiple reviews given by the same user exactly at the same time

## Business Question 1:

What are the 3 most important aspects focused in bad reviews?

# Dataset Preparation

- User Id over Profile Name
- Removed Null entries
- Removed repeated observations
- Removed entries where number of people who found the review helpful was greater than total surveyed
- Bad Ratings - 1 & 2 Ratings



# Model Analysis - Topic Modelling

- Topic models provide a simple way to extract topics and themes from large volumes of unlabeled text.
- Latent Dirichlet Allocation(LDA)
- Unsupervised Learning Technique
- Mallet LDA package
- MALLET is a Java-based package for statistical natural language processing, document classification, clustering, topic modeling.
- Efficient Parallel Processing



# Topic Modelling on Bad Reviews

Topic Modelling on text of bad reviews

```
```{r}
lower_text = tolower(baddf$Text)
ctext = Corpus(VectorSource(lower_text))
rm(data)

mallet.instances <- mallet.import(as.character(seq(1:length(lower_text))),
                                lower_text,
                                "/C:/Amrita/HardDisk/SpringTerm2018/MachineLearning/stopWords.txt")

topic.model <- MalletLDA(num.topics=3)
topic.model$loadDocuments(mallet.instances)
topic.model$setAlphaOptimization(20, 100) # optimise parameters after every 20 iterations which will be preceded by 100 burnin
topic.model$train(1000) # train the model
topic.model$maximize(10)
doc.topics <- mallet.doc.topics(topic.model, smoothed=T, normalized=T)
topic.words <- mallet.topic.words(topic.model, smoothed=T, normalized=T)
topics.labels <- rep("",3)
for (topic in 1:3) topics.labels[topic] <- paste(mallet.top.words(topic.model, topic.words[topic,], num.top.words=80)$words,
collapse=" ")
topics.labels
```
```

# Common Reasons of Complaints - Bad Reviews

- Health effects of defective food products.
- Misleading/deceptive products, shipping, packaging services, damaged and expired products.
- Taste, flavor and smell aspects of the product.

# Insights / Suggestions

- Scrutinize the vendors who are shipping defective products to improve customer experience
- Improve the delivery distribution channel and also analyze the underlying root causes of why the goods are damaged
- Packaging standards should also be improved for perishable/fragile items

## Business Question 2:

What are the most important characteristics of a Helpful Review?

# Why?

- 85% of customers trust online reviews as much as a personal recommendation
- Hence measuring helpfulness goes a long way in developing the trust of a customer
- Idea behind the question is to develop a model which would identify if a review would be helpful or not

# Dataset Preparation...

- Helpfulness Score - Metric to measure the helpfulness of each review
- Helpfulness Score = 
$$\frac{\text{Helpful}}{\text{Helpful} + \text{Unhelpful}}$$
- Ranges from 0 to 1.
- Created a Binary Variable with threshold as 0.5
- Reviews which didn't have a helpfulness voting
  - Removed those observations
  - Created a new category by considering these reviews as neutral

# Features

- Word Count
- Sentence Count
- Rating
- Sentiment Score
- Readability Score
- Similarity Score

# Sentiment Score

- Analyze whether strong sentiments make a review helpful or not
- SentimentR package
- Calculated Polarity Score for each review which ranges from -1 to +1

```
...{r}  
library(sentimentr)  
sentimentdf <- with(nonspamdf, sentiment_by(get_sentences(Text), list(Id)))  
write.csv(sentimentdf, file = "sentimentdf.csv", row.names=TRUE, na="")  
...
```



# Readability Index

- The index estimates the years of formal education needed to understand text on a first reading
- Readability package
- Gunning Fog Index per review
- Index ranges from 1-100.(Hard-Easy)
- Reviews with high index are easier to read than reviews with low

```
library(readability)
readable <- with(nonspamdf, readability(Text, list(Id)))
write.csv(readable, file = "readability.csv", row.names=TRUE, na="")
```

# Similarity Score

- Lexical similarity which determines how close a review is to the already identified helpful features
- Extracted top 200 words that occur the most in a helpful review
- How helpful is our review based on comparing each review to these extracted features

# Similarity Score...

```
# Taking out the helpful reviews
helpfulreviews=filtered[filtered$HelpfulnessScore_bin==1,]
summary(helpfulreviews$Text)

# Top 200 words
lower_text =tolower(helpfulreviews$Text)
ctext = Corpus(VectorSource(lower_text))
ctext_nopunc_nonum = tm_map(ctext, removeNumbers)
ctext_nopunc = tm_map(ctext_nopunc_nonum, removePunctuation)
ctext_nopunc_nonum_nostop = tm_map(ctext_nopunc, removeWords, c("shall","us","unto","will","just","nothing","can",
"much","dont","didnt","doesnt","never", "upon","also","let","even","now","yet", "therefore","may","away","since","nothing",
stopwords("english")))

tdm2 = TermDocumentMatrix(ctext_nopunc_nonum_nostop,control=list(wordLengths=c(4, 15),
                        bounds = list(global = c(50,Inf))))

tdm3 = as.matrix(tdm2)
wordcount = sort(rowSums(tdm3),decreasing=TRUE)
tdm_names = names(wordcount)[1:200]
wordcloud(tdm_names,wordcount)

# Similarity Calculation
m = length(filtered$Text) # No of sentences in input
text=filtered$Text
jaccard = matrix(0,m,1) #Store match index
b = tdm_names ; bb = unlist(strsplit(b," "))
for (i in 1:m) {
  a = text[i]; aa = unlist(strsplit(a," "))
  jaccard[i] = length(intersect(aa,bb))/
    length(union(aa,bb))
}
filtered$SimilarityScore=jaccard
```

## Approach 1: Binary Classification

| Model               | Accuracy | Precision | Recall | F-1 Score |
|---------------------|----------|-----------|--------|-----------|
| LDA                 | 76.9%    | 90.6%     | 81.6%  | 85.8%     |
| Logistic Regression | 77.3%    | 93.2%     | 80.6%  | 86.4%     |
| KNN                 | 69.6%    | 80.1%     | 80.7%  | 80.4%     |
| Naive Bayes         | 76.3%    | 88.2%     | 82.4%  | 85.2%     |
| XGboost             | 77.7%    | 95.6%     | 79.6%  | 86.9%     |

# Binary Classification Results

```
Call:
glm(formula = y_train ~ Gunning_Fog_Index + word_count + ave_sentiment +
     Score + SimiliarityScore, family = binomial(link = "logit"),
     data = x_train)
```

Deviance Residuals:

| Min     | 1Q     | Median | 3Q     | Max    |
|---------|--------|--------|--------|--------|
| -2.4436 | 0.4391 | 0.5464 | 0.6166 | 1.3949 |

Coefficients:

|                   | Estimate   | Std. Error | z value | Pr(> z ) |     |
|-------------------|------------|------------|---------|----------|-----|
| (Intercept)       | -0.9332187 | 0.0289627  | -32.221 | < 2e-16  | *** |
| Gunning_Fog_Index | 0.0027711  | 0.0023980  | 1.156   | 0.248    |     |
| word_count        | 0.0001410  | 0.0001264  | 1.115   | 0.265    |     |
| ave_sentiment     | -0.1937607 | 0.0408331  | -4.745  | 2.08e-06 | *** |
| Score             | 0.4877383  | 0.0052698  | 92.554  | < 2e-16  | *** |
| SimiliarityScore  | 11.3652526 | 0.5780690  | 19.661  | < 2e-16  | *** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 127577 on 119828 degrees of freedom  
Residual deviance: 115018 on 119823 degrees of freedom  
AIC: 115030

Number of Fisher Scoring iterations: 4

## Approach 2: Multinomial Classification Prep

- Reviews which didn't have a helpfulness voting are categorized in the neutral category
- 0 - Not Helpful
- 1 - Helpful
- 2 - Neutral
- Same features as binomial classification

# Multinomial Classification

| Model                  | Accuracy | Precision | Recall | F-1 Score |
|------------------------|----------|-----------|--------|-----------|
| LDA                    | 48.6%    | 27.9%     | 49.2%  | 35.6%     |
| KNN                    | 44.1%    | 44.1%     | 44.2%  | 44.2%     |
| Multinomial Regression | 49.0%    | 30.2%     | 48.6%  | 37.3%     |
| Naive Bayes            | 48.4%    | 20.5%     | 50.7%  | 29.2%     |
| XGBoost                | 50.2%    | 46.3%     | 40.7%  | 40.9%     |

# Insights

- Surprisingly, the number of words and sentences used in a review doesn't have an impact on the degree of helpfulness
- Reviews with strong negative sentiments are more helpful as compared to strong positive sentiments.
- Reviews containing more words which describe the product and packaging attributes and provide suggestions to the customers are more helpful



Thank you